



Aaditya Bhilegaonkar

📞 [6028042664](tel:6028042664) ✉ [Email](#) 🐙 [Github](#) in [Linkedin](#) 📝 [Medium Blog](#)

SUMMARY

Data Engineer with **2+ years** building production ETL pipelines and warehouses at scale. Skilled in Python, SQL, distributed systems, and AI-augmented workflows; active open source contributor across **Astronomer**, **Tesla**, **Kestra**, and **Apache**.

WORK EXPERIENCE

Data Engineering Intern

University of Phoenix

May 2025 - August 2025

- Accelerated **analytics-ready** data delivery by **30%** across **10+** sources by optimizing **Python/SQL** pipeline execution.
- Improved data availability by automating **Redshift** cleanup jobs using **Airflow**, saving **20+ engineering hours** per month.
- Redesigned validation logic by **decoupling** it from transformation layer, improving pipeline **performance** by **20%**.
- Reduced **Redshift** Spectrum query costs by **70%** on **100–200GB** weekly datasets by optimizing **S3** partition pruning.

Software Developer - Data

Amdocs

June 2022 - July 2024

- Restored billing accuracy for **200M+** subscribers, preventing **\$100K+/month** in miscalculations by fixing pipeline defect.
- Resolved **10M+** row data corruption across Oracle tables in under **7 hours** via **SQL** log analysis and schema impact assessment.
- Improved batch job runtime by **18%** on **500–800GB** telecom datasets by rewriting legacy SQL and optimizing execution logic.
- Implemented **data validation** across **5+** pipelines processing **100GB+/day**, reducing downstream errors for analytics teams.

EDUCATION

Arizona State University, Tempe, USA

May 2026

Master of Science in Computer Science

Pune Institute of Computer Technology, Pune, India

April 2019 - May 2022

Bachelor of Engineering in Electronics and Telecommunications

PROJECTS

Agentic Data Engineering Platform | Google ADK, BigQuery

- Cut ETL pipeline development time from **5+ hours** to minutes by building a **Google ADK** multi-agent system that automates data preparation, job scheduling, and **BigQuery** warehouse loading end-to-end.

Data Lineage Analysis Tool | Python, dbt, Airflow

- Cut breaking-change impact assessment from hours to seconds by engineering a **BFS** lineage tracer spanning **10+** file types (**dbt**, **Airflow**, PySpark, notebooks) with **5-stage** auto-generated artifacts per change.

LLM-Augmented Metadata Pipeline | PySpark, AWS Glue, LLaMA

- Reduced ingestion time by **20%** and analyst query turnaround by **30%** by building a **PySpark + AWS Glue** Medallion ETL pipeline with **LLaMA** trained on table metadata to auto-generate SQL.

Claude Code Developer Tools | TypeScript, Node.js

- Reduced **Claude Code** API token overhead by **80–300 tokens/request** by shipping a transparent HTTP proxy (**52 tests**, npm v0.3.3) that strips unused tool schemas and prunes conversation history on the wire.

Open Source Contributions | Python, Go, Java

- Contributed **7+ PRs** across **5 repos** spanning Python, Go, and Java: reduced dbt-Airflow graph complexity **11→9**, extended Snowflake profile mappings, improved CLI startup UX, and documented fleet-telemetry partition strategy.

RAG Retrieval Optimization | Milvus, Python

- Improved **Recall@5** by **14%** and search latency by **35%** by benchmarking hybrid chunking strategies across **3** retrieval modalities using **Milvus** with nDCG@10/Recall@5 evaluation pipelines.

TECHNICAL SKILLS

Programming Languages	: Python, TypeScript, Unix/Linux Scripting, SQL.
Cloud & Infrastructure	: AWS (S3, IAM, SQS), Docker, Redshift Spectrum.
Distributed Systems	: Spark, Kafka, Hive, Hadoop, Flink.
Data platforms & Warehousing	: BigQuery, Snowflake, Redshift, Dimensional Modeling, Star Schema, OLAP/OLTP.
Data Integration & Orchestration	: Airflow, AWS Glue, Fivetran, Tableau, dbt, ETL/ELT.
DevOps & Monitoring	: AWS CloudWatch, Jenkins, Git, CI/CD.
Agentic AI & LLM Systems	: Google ADK, MCP Framework, Milvus, RAG Architecture.